

# Lock-free programming: MPSC queue

Donnerstag, 6. Februar 2025 09:16

Author: Damir Ljubic  
email: damirlj@yahoo.com

## Introduction

I'm keep saying - beside the software architecture and design, as most fundamental, entry point of every software development, the alphabet of the programming that everyone should master: is **asynchronous programming**.

I would even dare to say - the most of the time in the implementation phase I'm spending dealing with all kind of "*asynchronicity*", and how to solve these challenges within the given constraints: available hardware, toolchain, framework, even language.

From the language perspective, the C++ influenced by the other languages and libraries made a significant effort to provide the mechanisms for concurrency as part of the standard library itself, starting with a thread abstraction, synchronization primitives and a new memory model - all introduced with C++11.

These enabled us to address concurrency topics from different, non-conventional point of view: writing **lock-free** (even wait-free) code, for potentially overcoming the lock-based approach main issues:

→ The performance and dead-locks.

Potentially because the operations on atomic types (atomic operations) are notably more expensive then the same operations on the non-atomic types, and may even involve the locking mechanism.

There is also *Cashe coherence* involved - synchronization between the cores hosting the threads synchronized around the same atomics, etc. - but these are usually the minor issues on contemporary platforms.

The main pre-warning for lock-free programming is that it's difficult to write it correctly.

→ *What does it mean?*

With lock-based programming - we have a **strong thread-safety guaranties**, that our Jailkeeper will not left any between inmates shared rooms without monitoring, as long as there is possibility for the prison brake: as long as there is possibility for the race condition and undefined behavior.

With lock-free programming, we are juggling with the trade-safety guaranties, being much more "*relaxed*".

We are talking about the **memory ordering** and the **visibility** of the same code in multithread context - non-atomic operations around the atomic variables used for synchronizing the multiple threads.

→ More on that: [https://github.com/damirlj/modern\\_cpp\\_tutorials/blob/main/docs/memory/Lock-free%20programming%2C%20part%201.pdf](https://github.com/damirlj/modern_cpp_tutorials/blob/main/docs/memory/Lock-free%20programming%2C%20part%201.pdf)

This allows us more flexibility - especially designing the lock-free structures so that they reflect certain use-cases, certain scenarios where we conditionally can lowering these expectations on a thread-safety: not being so rigid Jailkeeper, but rather the one who closes one eye on race conditions - and still maintains the order, having all prisoners on count.

Enough with the prisons analogies.

Let's examine some concrete implementation.

## MPSC - Multiple Producers Single Consumer lock-free queue

Relaxing these thread-safety guaranties we simplify the implementation of the lock-free structures, that will be tailored to work correctly - in line with expectations only under certain circumstances.

The most simple one, in terms of the lock-free queue would be Single Producer Single Consumer implementation.

In the *lock-based* implementation, for bounded queue, we would have a `std::array` as memory storage and two semaphores (`std::counting_semaphore`), where the internal counter for the writing semaphore is initially set to the number of empty slots in memory storage, while the read semaphore will be initially set to 0 - as indication of the actual number of elements.

The producer acquires the write semaphore - that will block if there is no empty slots, otherwise the producer will write in the first empty slot and notify the consumer - by releasing the read semaphore.

The consumer will acquire the read semaphore - and wait if the internal counter is 0.

Otherwise, it will proceed - pop the queue and release the write semaphore.

→ Possible implementation: [https://github.com/damirlj/modern\\_cpp\\_tutorials/tree/main/src/ring%20buffer](https://github.com/damirlj/modern_cpp_tutorials/tree/main/src/ring%20buffer)

For *lock-free* version of the queue, we would use `std::atomic<std::size_t>` - to represent the write/read index

```
alignas(64) std::array<T, N> data_;  
alignas(64) std::atomic<std::size_t> head_ {0};  
alignas(64) std::atomic<std::size_t> tail_ {0};
```

Assuming that producer will ever call *push* method, for SPSC implementation we could write something like

```
template <typename U>
requires std::convertible_to<U, value_type>
bool try_push(U&& u) noexcept(std::is_nothrow_constructible_v<U>)
{
    const auto tail = tail_.load(std::memory_order_relaxed); // maintain by the single producer
    const auto head = head_.load(std::memory_order_acquire); // maintain by the single consumer
    if ( ((tail + 1) & (N - 1)) == head) return false; // queue is full

    data_[tail] = std::forward<U>(u);
    tail_.store((tail + 1) & (N - 1), std::memory_order_release); // signal write event to consumer

    return true;
}
```

For MPSC scenario, we need to take into consideration that multiple producers can advance the tail index simultaneously, so we need to check the "*full buffer*" condition repeatedly

```
template <typename U>
requires std::convertible_to<U, value_type>
bool push(U&& u) noexcept(std::is_nothrow_constructible_v<U>)
{
    // expected value - otherwise, another producer modifies it
    auto tail = tail_.load(std::memory_order_relaxed);
    while (is_full() || not tail_.compare_exchange_weak(tail, (tail + 1) & (N - 1), std::memory_order_release));

    data_[tail] = std::forward<U>(u);

    return true;
}
```

The consumer part remains the same, in terms of the *dequeuing*, since we have a single consumer  
We have a private utility method

```
template<typename Func, typename...Args>
requires std::invocable<Func, std::size_t, Args...> && std::is_same_v<bool, std::invoke_result_t<Func,
std::size_t, Args...>>
inline std::optional<value_type> pop(Func&& func, Args&&...args)
noexcept(std::is_nothrow_move_constructible_v<value_type>)
{
    const auto head = head_.load(std::memory_order_relaxed);

    if(not std::invoke(std::forward<Func>(func), head, std::forward<Args>(args)...)) return {};

    auto data = std::optional<value_type>(std::move(data_[head]));

    head_.store((head + 1) & (N - 1), std::memory_order_release);

    return data;
}
```

We obtain the head index maintained within the single, consumer-thread (relaxed), and in case that the given condition is fulfilled - the data in buffer on the given slot will be moved, and afterwards the head index will be advanced - with *release memory barrier*: to synchronize with producers happens-before relationship, as they are checking "*full buffer*" condition being fulfilled.

We have, similar as with enqueueing mechanism, different variants of *pop* method

*try\_pop()* : try to dequeue the element from the queue, if the queue is not empty

```
auto try_pop() -> std::optional<value_type>
{
    return pop([this](std::size_t head) {return not is_empty(head); });
}
```

For the waiting dequeue operations, until the event is moved, or the stop token is signaled

```
auto pop_wait(const std::atomic_flag& stop) ->std::optional<value_type>
{
    return pop([&stop, this](std::size_t head)
    {
        while (is_empty(head)) // wait until is non-empty, or stop is signaled
    });
}
```

```

        {
            if (stop.test(std::memory_order_relaxed)) return false;
            std::this_thread::yield();
        }
        return true;
    });
}

```

The same as previous, but with additional exit criteria: timeout is expired

```

auto pop_wait_for(const std::atomic_flag& stop, std::chrono::milliseconds timeout) ->std::optional<value_type>
{
    return pop([&stop, timeout, this](std::size_t head)
    {
        using namespace std::chrono;
        auto start = steady_clock::now();

        while (is_empty(head)) // wait until is non-empty, stop is signaled, or timeout expired
        {
            if (stop.test(std::memory_order_relaxed)) return false;
            if (duration_cast<milliseconds>(steady_clock::now() - start) > timeout) return false;
            std::this_thread::yield();
        }
        return true;
    });
}

```

For more details, check the entire implementation

<Compiler Explorer>: <https://godbolt.org/z/8T4vqrWoE>