

Lock-free programming: MPSC queue

Donnerstag, 6. Februar 2025 09:16

Author: Damir Ljubic
email: damirlj@yahoo.com

Introduction

I'm keep saying - beside the software architecture and design, as most fundamental, entry point of every software development, the alphabet of the programming that everyone should master: is **asynchronous programming**.

I would even dare to say - the most of the time in the implementation phase I'm spending dealing with all kind of "*asynchronicity*", and how to solve these challenges within the given constraints: available hardware, toolchain, framework, even language.

From the language perspective, the C++ influenced by the other languages and libraries made a significant effort to provide the mechanisms for concurrency as part of the standard library itself, starting with a thread abstraction, synchronization primitives and a new memory model - all introduced with C++11.

These enabled us to address concurrency topics from different, non-conventional point of view: writing **lock-free** (even wait-free) code, for potentially overcoming the lock-based approach main issues:

→ The performance and dead-locks.

Potentially because the operations on atomic types (atomic operations) are notably more expensive then the same operations on the non-atomic types, and may even involve the locking mechanism.

There is also *Cashe coherence* involved - synchronization between the cores hosting the threads synchronized around the same atomics, etc. - but these are usually the minor issues on contemporary platforms.

The main pre-warning for lock-free programming is that it's difficult to write it correctly.

→ *What does it mean?*

With lock-based programming - we have a **strong thread-safety guaranties**, that our Jailkeeper will not left any between inmates shared rooms without monitoring, as long as there is possibility for the prison brake: as long as there is possibility for the race condition and undefined behavior.

With lock-free programming, we are juggling with the trade-safety guaranties, being much more "*relaxed*".

We are talking about the **memory ordering** and the **visibility** of the same code in multithread context - non-atomic operations around the atomic variables used for synchronizing the multiple threads.

→ More on that: https://github.com/damirlj/modern_cpp_tutorials/blob/main/docs/memory/Lock-free%20programming%2C%20part%201.pdf

This allows us more flexibility - especially designing the lock-free structures so that they reflect certain use-cases, certain scenarios where we conditionally can lowering these expectations on a thread-safety: not being so rigid Jailkeeper, but rather the one who closes one eye on race conditions - and still maintains the order, having all prisoners on count.

Enough with the prisons analogies.

Let's examine some concrete implementation.

MPSC - Multiple Producers Single Consumer lock-free queue

Relaxing these thread-safety guaranties we simplify the implementation of the lock-free structures, that will be tailored to work correctly - in line with expectations only under certain circumstances.

The most simple one, in terms of the lock-free queue would be Single Producer Single Consumer implementation.

In the *lock-based* implementation, for bounded queue, we would have a `std::array` as memory storage and two semaphores (`std::counting_semaphore`), where the internal counter for the writing semaphore is initially set to the number of empty slots in memory storage, while the read semaphore will be initially set to 0 - as indication of the actual number of elements.

The producer acquires the write semaphore - that will block if there is no empty slots, otherwise the producer will write in the first empty slot and notify the consumer - by releasing the read semaphore.

The consumer will acquire the read semaphore - and wait if the internal counter is 0.

Otherwise, it will proceed - pop the queue and release the write semaphore.

→ Possible implementation: https://github.com/damirlj/modern_cpp_tutorials/tree/main/src/ring%20buffer

For *lock-free* version of the queue, we would use `std::atomic<std::size_t>` - to represent the write/read index

```
alignas(64) std::array<T, N> data_ ;
alignas(64) std::atomic<std::size_t> head_ {0};
alignas(64) std::atomic<std::size_t> tail_ {0};
```

Assuming that producer will ever call *push* method, for SPSC implementation we could write something like

```
template <typename U>
requires std::convertible_to<U, value_type>
bool try_push(U&& u) noexcept(std::is_nothrow_constructible_v<U>)
{
    const auto tail = tail_.load(std::memory_order_relaxed); // maintain by the single producer
    const auto head = head_.load(std::memory_order_acquire); // maintain by the single consumer
    if ( ((tail + 1) & (N - 1)) == head) return false; // queue is full

    data_[tail] = std::forward<U>(u);
    tail_.store((tail + 1) & (N - 1), std::memory_order_release); // signal write event to consumer

    return true;
}
```

For MPSC scenario, we need to take into consideration that multiple producers can advance the tail index simultaneously, so we need to check the "full buffer" condition repeatedly

```
template <typename U>
requires std::convertible_to<U, value_type>
void push(U&& u) noexcept(std::is_nothrow_constructible_v<U>)
{
    // expected value - otherwise, another producer modifies it
    auto tail = tail_.load(std::memory_order_relaxed);
    while (is_full() || not tail_.compare_exchange_weak( tail,
                                                         inc(tail),
                                                         std::memory_order_acq_rel,
                                                         std::memory_order_relaxed
                                                         ));

    data_[tail] = std::forward<U>(u);
}
```

The consumer part remains the same, in terms of the *dequeuing*, since we have a single consumer
We have a private utility method

```
template<typename Func, typename...Args>
requires std::invocable<Func, Args...> && std::lissome<bool, std::invoke_result_t<Func, Args...>>
inline std::optional<value_type> pop(Func&& func, Args&&...args)
noexcept(std::is_nothrow_move_constructible_v<value_type>)
{
    if (not std::invoke(std::forward<Func>(func), head, std::forward<Args>(args)...)) return {};

    const auto head = head_.load(std::memory_order_relaxed);
    auto data = std::optional<value_type>(std::move(data_[head]));

    head_.store((head + 1) & (N - 1), std::memory_order_release);

    return data;
}
```

We obtain the head index maintained within the single, consumer-thread (relaxed), and in case that the given condition is fulfilled - the data in buffer on the given slot will be moved, and afterwards the head index will be advanced - with *release memory barrier*: to synchronize with producers happens-before relationship, as they are checking "full buffer" condition being fulfilled.

We have, similar as with enqueueing mechanism, different variants of *pop* method

try_pop() : try to dequeue the element from the queue, if the queue is not empty

```
auto try_pop() -> std::optional<value_type>
{
    return pop([this]() {return not is_empty(); });
}
```

For the waiting dequeue operations, until the event is moved, or the stop token is signaled

```
auto pop_wait(const std::atomic_flag& stop) -> std::optional<value_type>
{
    return pop([&stop, this]()
```

```

    {
        while (is_empty()) // wait until is non-empty, or stop is signaled
        {
            if (stop.test(std::memory_order_relaxed)) return false;
            std::this_thread::yield();
        }
        return true;
    });
}

```

The same as previous, but with additional exit criteria: timeout is expired

```

auto pop_wait_for(const std::atomic_flag& stop, std::chrono::milliseconds timeout) -> std::optional<value_type>
{
    return pop([&stop, timeout, this]()
    {
        using namespace std::chrono;
        auto start = steady_clock::now();

        while (is_empty()) // wait until is non-empty, stop is signaled, or timeout expired
        {
            if (stop.test(std::memory_order_relaxed)) return false;
            if (duration_cast<milliseconds>(steady_clock::now() - start) > timeout) return false;

            std::this_thread::yield();
        }
        return true;
    });
}

```

For more details, check the entire implementation

<Compiler Explorer>: <https://godbolt.org/z/G1hxheoEd>

Addendum

The MPMC - Multi Producers Multi Consumers queue implementation

The additional constraints need to be imposed now on the dequeuing side, since we are extending the usage on the multiple consumers that may simultaneously access the same data structure: attempting to drain the queue
The implementation becomes quite symmetric to the handling of multiple concurrent enqueueing operations.

```

std::optional<value_type> pop(const std::atomic_flag& stop)
noexcept(std::is_nothrow_move_constructible_v<value_type>)
{
    for(;;)
    {
        auto head = head_.load(std::memory_order_relaxed);

        if (not is_empty(head) && head_.compare_exchange_strong(
            head,
            inc(head),
            std::memory_order_acq_rel,
            std::memory_order_relaxed))
        {
            return std::optional<value_type>(std::move(data_[head]));
        }

        if (stop.test(std::memory_order_relaxed)) break; // forcibly exit

        std::this_thread::yield();
    }

    return {};
}

```


<Compiler Explorer>: <https://godbolt.org/z/bvG4qT7KT>

Well - almost.

With this approach - the multiple consumers still may read out of sequence.

This is why we need to add additional constraint - **FIFO compliant sequencing**.

The producer threads will write $index + 1$ into slot - to indicate that the slot is written, while the consumers will write $index + MASK + 1$ - to indicate that the slot is empty



<Final version>: <https://godbolt.org/z/dhMhdsEhv>

Links

[Dmitry Vyukov]: https://github.com/craflin/LockFreeQueue/blob/master/mpmc_bounded_queue.h